

Releases

Table of contents

1 Download.....	2
2 Release Notes.....	2
3 News.....	2
3.1 9 October, 2007: release 0.14.2 available	2
3.2 4 September, 2007: release 0.14.1 available	2
3.3 25 July, 2007: release 0.13.1 available.....	2
3.4 8 June, 2006: FAQ added to Wiki	3

1. Download

Releases may be downloaded from Apache mirrors.

[Download a release now!](#)

On the mirror, all recent releases are available, but are not guaranteed to be stable. For stable releases, look in the stable directory.

2. Release Notes

Release notes for Hadoop releases are available in Jira.

[Browse release notes now!](#)

3. News

3.1. 9 October, 2007: release 0.14.2 available

This release fixes critical bugs in release 0.14.1.

3.2. 4 September, 2007: release 0.14.1 available

New features in release 0.14 include:

- Better checksums in HDFS. Checksums are no longer stored in parallel HDFS files, but are stored directly by datanodes alongside blocks. This is more efficient for the namenode and also improves data integrity.
- Pipes: A C++ API for MapReduce
- Eclipse Plugin, including HDFS browsing, job monitoring, etc.
- File modification times in HDFS.

There are many other improvements, bug fixes, optimizations and new features. Performance and reliability are better than ever.

Note:

When upgrading an existing HDFS filesystem to this release, you should first start HDFS with 'bin/start-dfs.sh -upgrade'. See the [Hadoop 0.14 Upgrade](#) page for details.

3.3. 25 July, 2007: release 0.13.1 available

In addition to new features, release 0.13 contains significant performance and reliability

improvements. For example, MapReduce sort performance on large clusters has [doubled](#).

Note:

When upgrading an existing HDFS filesystem to this release, you should first start HDFS with `bin/start-dfs.sh -upgrade`. See the [Hadoop Upgrade](#) page for details.

3.4. 8 June, 2006: FAQ added to Wiki

Hadoop now has a [FAQ](#). Please help make this more complete!